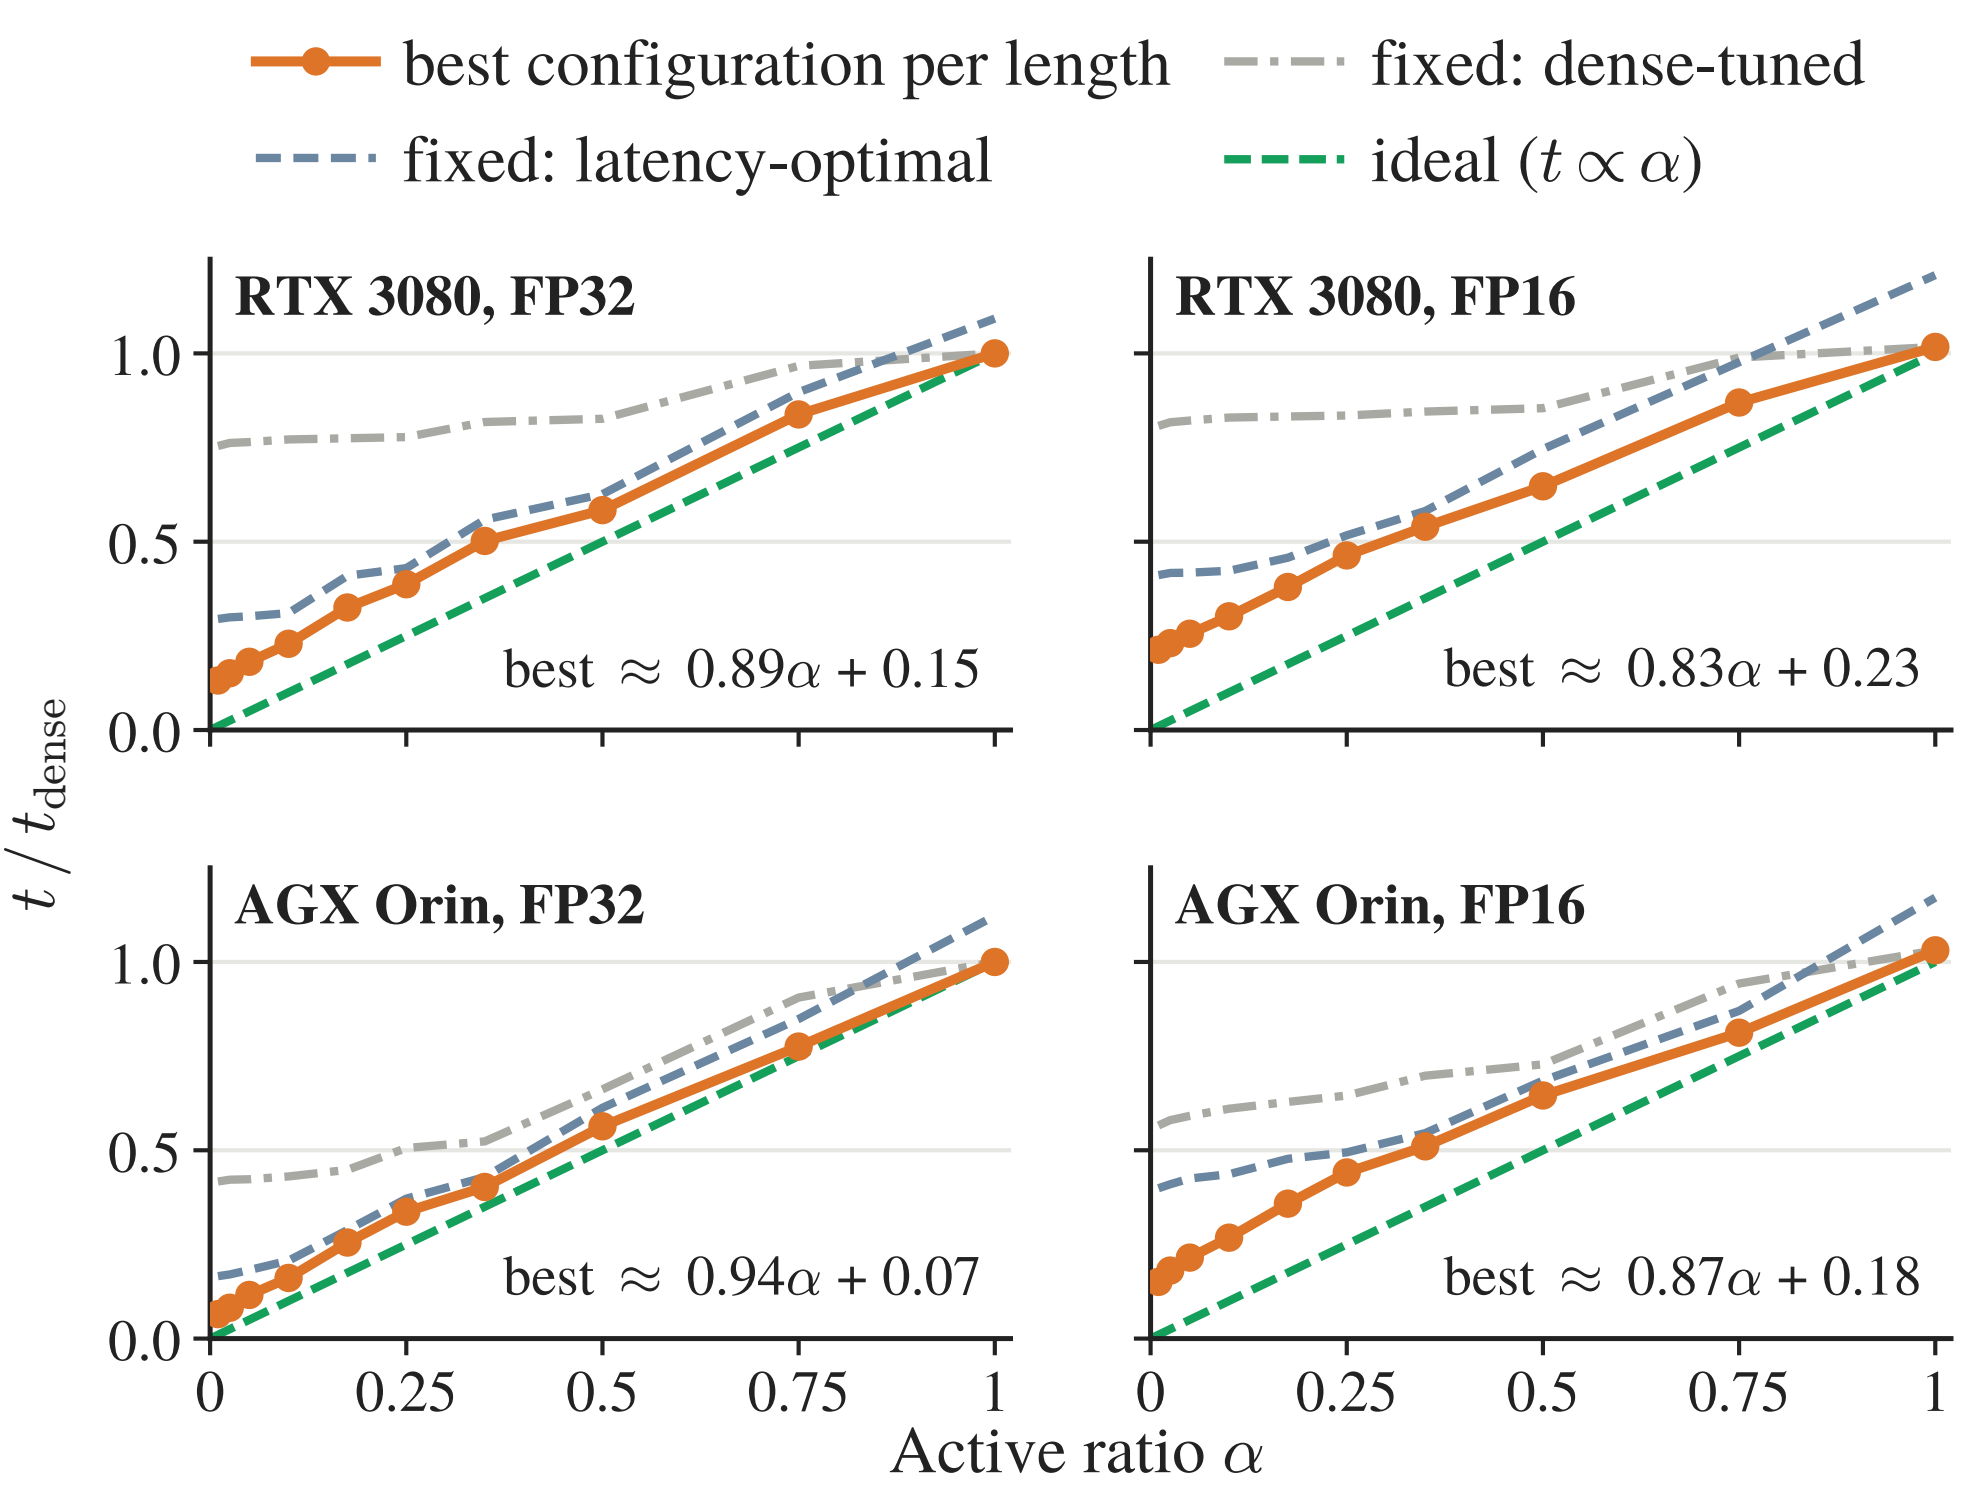
best configuration per length
fixed: dense-tuned
fixed: latency-optimal
ideal ($t \propto \alpha$)
RTX 3080, FP32
best $\approx 0.89\alpha + 0.15$
RTX 3080, FP16
best $\approx 0.83\alpha + 0.23$
AGX Orin, FP32
best $\approx 0.94\alpha + 0.07$
AGX Orin, FP16
best $\approx 0.87\alpha + 0.18$
t / t_{dense}
Active ratio α
0.0
0.5
1.0
0
0.25
0.5
0.75
1