

WISEConv: Worklist-driven Masked Convolution for Onboard High-Speed Robotic Perception

Anonymous Authors

Abstract—Onboard robotic perception relies on 2D convolutional networks that must meet hard latency constraints under the limited compute of embedded platforms. Many high-speed vision systems exploit spatial sparsity by producing an active mask and updating only the affected convolution outputs. However, existing approaches face a fundamental trade-off: tile skipping preserves the dense pipeline’s throughput but computes many unneeded outputs, while gather-scatter computes only needed outputs but sacrifices throughput. Neither converts high sparsity into proportional latency reduction. We present WISEConv, a masked convolution operator that achieves both high useful-compute ratio and near-dense throughput by fusing position-level selectivity into the dense pipeline itself. WISEConv first scans the mask to build an ordered worklist of active positions, then drives the standard convolution datapath using only those positions. On Jetson Xavier NX and Orin, across three mask sources and three robotic perception tasks, WISEConv achieves near-linear latency scaling with the active ratio, outperforming state-of-the-art baselines while preserving task accuracy.

Index Terms—Computer architecture for robotic and automation, software, deep learning methods, sparse convolution, embedded perception.

I. INTRODUCTION

ONBOARD robots run convolutional perception in the control loop for navigation, manipulation, and closed-loop control [1]–[3]. On embedded platforms the available compute is far below server-class GPUs, yet perception latency is a hard real-time constraint because it sits directly on the control path. Dense convolution recomputes every spatial position, and on these platforms it often cannot reach the update frequency the loop demands. Many high-speed vision systems on UAVs, mobile robots, and embedded devices [4]–[7] therefore concentrate compute on active regions via *masked convolution*: the input stays a dense image-grid tensor, but only the outputs at positions marked active are updated.

The standard dense convolution pipeline partitions the output grid into *tiles* (contiguous spatial blocks) and for each tile reads the required input patch by coordinate, computes the convolution, and writes the result back. This coordinate-driven, tile-based execution achieves high throughput because tiles access spatially contiguous data with regular, predictable memory patterns. Existing masked-convolution systems modify or abandon this pipeline along two paths.

Tile skipping [4], [6], [8], [9] stays inside the dense pipeline: it tests each tile and skips it if every position within is inactive. Throughput is preserved for non-skipped tiles, but a single active position forces the entire tile to execute, so many computed outputs are unneeded. *Gather-scatter* [7], [10] abandons the pipeline entirely: it extracts active positions exactly, gathers

their input neighborhoods, computes, and scatters results back. Every computed output is needed, but the extraction, irregular memory access, and scatter writeback collapse throughput.

We characterize this trade-off along two axes: *throughput* (operations per second, bounded by the dense pipeline) and *useful-compute ratio* (the fraction of executed operations that produce needed outputs). Tile skipping preserves throughput but sacrifices useful-compute ratio; gather-scatter achieves high useful-compute ratio but sacrifices throughput. Neither delivers both, so upstream sparsity does not translate into proportional latency reduction. In our evaluation on Jetson Xavier NX and Orin, existing methods [4], [6], [7], [10], [11] reduce the active-output ratio to roughly XX–YY%; yet tile-skipping methods achieve only AA–BB% useful-compute ratio, while gather-scatter methods reach only CC% of tile-skipping throughput. Convolution latency is reduced by only DD–EE%.

Our key insight is that the position-level selectivity of gather-scatter can be fused into the dense pipeline’s coordinate-driven execution without the separate data-movement stages. The fusion is straightforward in principle: because each tile already performs a coordinate-driven read-compute-write per position, replacing the tile’s coordinate source, from a contiguous block to a set of active positions, gives position-level selectivity while keeping the per-position execution path identical. This achieves high useful-compute ratio directly, and because the per-position datapath remains that of the dense pipeline, it opens the potential for dense-level throughput.

Realizing this potential requires addressing the challenge of two practical overheads. First, *coordinate discovery*: the kernel must know which positions are active before computation begins, yet scanning for them inside the compute kernel would stall the main datapath. Second, *locality*: if active positions are processed in arbitrary (instead of contiguous) order, memory access loses the spatial regularity that gives the dense pipeline its speed; a global sort would restore order but is itself expensive.

We present WISEConv (**W**orklist-driven **maSkEd** Convolution), a dense-grid masked convolution operator that resolves both overheads through a two-stage design. A lightweight construction stage scans the active mask and emits a *worklist*: the ordered coordinates of active output positions. The subsequent compute stage drives the dense pipeline’s coordinate-driven execution with the worklist as its coordinate source, computing only the active outputs. Both overheads are addressed by the construction stage’s tiled scan: it operates entirely in 2D mask space (no feature data touched), so coordinate discovery costs only XX% of the compute-stage time; and it compacts active

positions from each tile into contiguous worklist segments, so the compute stage accesses spatially nearby input data when processing consecutive worklist entries, preserving locality comparable to the dense traversal without any global sort. With construction overhead minimized and only active outputs computed, WISEConv achieves near-linear latency scaling: an active-output ratio of $x\%$ yields convolution latency close to $x\%$ of the dense baseline.

This paper makes the following contributions:

- We identify the useful-compute ratio vs. throughput trade-off as the fundamental bottleneck preventing masked convolution from delivering proportional speedup on embedded platforms [4], [6]–[10].
- We present WISEConv, a worklist-driven masked convolution operator that fuses position-level selectivity into the dense pipeline’s coordinate-driven execution, achieving both high useful-compute ratio and near-dense throughput.
- On Jetson Xavier NX and Orin, WISEConv achieves near-linear latency scaling with active-output ratio across three mask sources and three robotic perception tasks, outperforming dense cuDNN, tile skipping, and gather-scatter while preserving task accuracy.

II. BACKGROUND AND MOTIVATION

A. 2D Dense Convolution on Embedded Platforms

Convolutional neural networks (CNNs) are the computational backbone of onboard robotic visual perception. CNN-based depth estimation and optical flow [1], [2], [12] feed obstacle avoidance and path planning; object detection [13] and semantic segmentation [14] provide scene understanding for navigation; and pose estimation [3], [15] enables closed-loop manipulation. While vision transformers [16] have achieved strong accuracy on these tasks, their quadratic attention cost makes them impractical under the tight latency and power budgets of embedded platforms; CNN-based models remain the dominant deployed choice for real-time onboard visual inference. Across these models, convolutional layers operating on image-grid feature maps dominate the inference workload and sit directly on the control-critical path.

On embedded platforms (Jetson Xavier NX, Orin, edge TPUs), power and thermal budgets limit available compute to a fraction of server-class GPUs, yet the closed control loop imposes a hard latency ceiling: perception must complete within the sensor frame interval (typically tens of milliseconds) to feed downstream planning and control in time.

Consider a CNN with L convolution layers. We describe the operation of a single layer l ; all per-layer quantities (H_{out} , W_{out} , C_{in} , C_{out} , k , M , etc.) are implicitly indexed by l but we omit the superscript for clarity unless comparing across layers.

A 2D convolution layer transforms an input feature map $\mathbf{x} \in \mathbb{R}^{H_{in} \times W_{in} \times C_{in}}$ into an output $\mathbf{y} \in \mathbb{R}^{H_{out} \times W_{out} \times C_{out}}$ using a weight tensor $\mathbf{W} \in \mathbb{R}^{C_{out} \times C_{in} \times k \times k}$, where k is the kernel size and $k \times k$ is the receptive field size. The computation is defined per spatial position $p = (h, w)$ in the output: for each

p , let $\mathcal{N}(p)$ denote the k^2 input positions in the receptive field of p . The layer computes:

$$\mathbf{y}[p] = \sum_{i=1}^{k^2} \mathbf{W}_i \mathbf{x}[\mathcal{N}_i(p)], \quad (1)$$

where $\mathcal{N}_i(p)$ is the i -th position in the receptive field and $\mathbf{W}_i \in \mathbb{R}^{C_{out} \times C_{in}}$ is the corresponding weight slice. The output at each position is thus an accumulation of k^2 weighted contributions from the input.

To formalize the GPU execution model, we define the *output index grid* $G_{out} \in \mathbb{Z}^{H_{out} \times W_{out}}$ with entries $G_{out}[h, w] = h \cdot W_{out} + w$, assigning each output spatial position a unique row-major linear index. The GPU partitions this grid into $N = N_h \cdot N_w$ disjoint tiles of size $t_h \times t_w$, where $N_h = \lceil H_{out}/t_h \rceil$ and $N_w = \lceil W_{out}/t_w \rceil$. Tiles are enumerated in row-major order as $T_1, \dots, T_N$. The n -th tile occupies row $i_n = \lfloor (n-1)/N_w \rfloor$ and column $j_n = (n-1) \bmod N_w$ of the tile grid, and is defined as the row-major flattening of its slice:

$$T_n = \text{vec}(G_{out}[i_n t_h : (i_n + 1)t_h, j_n t_w : (j_n + 1)t_w]^\top), \quad (2)$$

yielding a vector of $t_h \cdot t_w$ linear indices. Tiles are processed in parallel; within each tile, consecutive entries correspond to spatially adjacent positions whose receptive fields overlap and lie in contiguous memory. This regularity enables cuDNN-class [17] throughput. We denote by $\mathcal{C}$ the set of positions actually computed by a given execution strategy. Dense convolution computes all $H_{out} \times W_{out}$ positions unconditionally, which on embedded platforms often exceeds the latency budget.

B. Masked Convolution for Onboard Perception

1) *Mask Sources*: Many onboard vision systems exploit the fact that only a subset of the input changes between inference calls or is relevant to the current decision. These systems identify an input active mask $M_{in} \in \{0, 1\}^{H_{in} \times W_{in}}$ marking positions where the input has changed or is task-relevant. An output position p must be recomputed whenever its receptive field contains an active input position, giving the output active mask $M \in \{0, 1\}^{H_{out} \times W_{out}}$:

$$M[p] = \begin{cases} 1 & \text{if } \exists r \in \mathcal{N}(p) \text{ s.t. } M_{in}[r] = 1, \\ 0 & \text{otherwise.} \end{cases} \quad (3)$$

In a multi-layer network, the output mask of layer l becomes the input mask of layer $l+1$, propagating sparsity through the model. Three families of mask sources produce the initial M_{in} :

a) *Event-camera increments*.: Event cameras [18] are bio-inspired sensors in which each pixel independently fires when its brightness change exceeds a threshold, generating data only where intensity varies. They offer microsecond temporal resolution, high dynamic range, and negligible motion blur, making them widely used on UAVs and high-speed platforms. Events accumulated over a temporal window are differenced to reveal which positions received new activity, producing a naturally sparse mask [4], [5].

b) *Temporal frame differences.*: Standard frame-based cameras on mobile robots and autonomous vehicles capture video at a fixed rate. Between consecutive frames most of the scene remains static, especially under moderate ego-motion. Temporal differencing methods [6], [11] compare consecutive frames (or motion-compensated histories) and identify positions where pixel differences exceed a threshold; only outputs affected by these positions need recomputation.

c) *Learned spatial gating.*: Rather than relying on temporal change, dynamic networks [7] train a lightweight gating module that examines the current input and predicts, per spatial region, whether computation is needed for the downstream task. Regions classified as low-value are masked out, adapting the compute budget to task-relevant content per frame.

2) *Execution Paradigms*: Given a mask M , the system must compute $y[p]$ only for positions where $M[p] = 1$. Existing methods follow one of two paradigms:

a) *Tile skipping.*: A tile T_n is executed if and only if it contains at least one active position; otherwise it is skipped entirely. The computed set is $\mathcal{C} = \bigcup_{\{n: \exists g \in T_n, M[g]=1\}} T_n$. Because non-skipped tiles run the unmodified dense pipeline, throughput is preserved. However, $|\mathcal{C}|$ can be much larger than $\|M\|_0$ when active positions are scattered across many tiles: a single active position forces its entire tile to execute. SBNNet [8], Skip-Convolutions [9], Ev-Conv [4], and DeltaCNN [6] follow this paradigm.

b) *Gather-scatter.*: Active positions are extracted exactly: the system gathers $\{x[\mathcal{N}(g)] : M[g] = 1\}$ into a contiguous buffer, computes Eq. (1) for each, and scatters results back to the output grid. The computed set $\mathcal{C}$ contains only active positions, so no computation is wasted. However, the gather, scatter, and irregular memory access patterns abandon the dense pipeline's regular execution, reducing throughput to a fraction of tile-based methods. DGNNet [7] and SAI [10] follow this paradigm.

C. The Gap: Sparsity Does Not Become Speedup

For a network with L convolution layers, let $\|M^l\|_0$ denote the number of active positions (nonzeros) in layer l 's output mask. We define layer l 's *active ratio* $\rho^l = \|M^l\|_0 / (H_{out}^l W_{out}^l)$ and the network-level active ratio $\hat{\rho} = \frac{1}{L} \sum_{l=1}^L \rho^l$, the average fraction of output positions that require computation across all layers. We characterize masked-convolution efficiency along two axes: *throughput*, the rate of convolution operations per second (bounded by the dense pipeline), and *useful-compute ratio*, the fraction of executed operations that produce needed outputs ($\|M\|_0 / |\mathcal{C}|$). Tile skipping preserves throughput but sacrifices useful-compute ratio; gather-scatter achieves high useful-compute ratio but sacrifices throughput. Neither delivers both, so upstream sparsity does not translate into proportional latency reduction. Figure 1 quantifies this on Jetson Orin running XX on YY dataset: with $\hat{\rho}$ in the range XX–YY%, tile-skipping methods achieve only AA–BB% useful-compute ratio and gather-scatter methods reach only CC% of tile-skipping throughput, reducing convolution latency by only DD–EE% rather than the ideal $\hat{\rho}$ -linear scaling. This gap motivates WISEConv.

[TODO] Active ratio $\hat{\rho}$ (x-axis) vs. actual speedup over dense (y-axis) on Jetson Orin. Three mask sources, two paradigms. Dashed line = ideal linear scaling.

Fig. 1. Active ratio vs. actual convolution speedup on Jetson Orin for the three mask sources. The dashed line marks ideal (linear) scaling. Tile skipping and gather-scatter both fall far short.

D. 3D Sparse Convolution: A Different Problem

3D sparse convolution (SpConv v2 [19], TorchSparse++ [20]) targets point-cloud workloads where input data is inherently irregular and the core challenge is imposing computational structure on scattered voxels. WISEConv faces a fundamentally different problem: the input resides on a regular dense grid with structured, stride-based access; the challenge is how to skip inactive outputs without leaving the dense pipeline's already-efficient execution. Applying 3D sparse convolution directly to this 2D scenario performs poorly because it requires building coordinate hash tables and neighbor maps with mandatory GPU synchronization, and the per-frame mask changes prevent amortizing this cost across layers. On Jetson Orin, SpConv v2 applied to a 2D masked convolution with $\rho^l = \text{XX}\%$ runs $\text{XX}\times$ slower than the dense baseline due to construction overhead alone (Section ??).

III. WISECONV DESIGN

A. Overview

WISEConv resolves the useful-compute ratio vs. throughput tension (Section II-C) through a two-stage design (Figure 2):

- 1) **Construction stage.** A lightweight kernel scans the input active mask M_{in} and evaluates Eq. (3) for each output position, emitting an ordered *worklist* $\mathbf{q} = (q_1, \dots, q_m)$ of $m = \|M\|_0$ active output indices drawn from G_{out} . The worklist preserves tile-local spatial locality (Section III-B).
- 2) **Compute stage.** A convolution kernel consumes $\mathbf{q}$ as its coordinate source. For each q_j , the kernel decodes the spatial coordinate, reads the receptive-field input $x[\mathcal{N}(q_j)]$, and accumulates Eq. (1) identically to the dense pipeline (Section III-C).

The computed set $\mathcal{C} = \{q_1, \dots, q_m\}$ contains only active positions (high useful-compute ratio), while the per-position datapath remains that of the dense pipeline (near-dense throughput). The construction stage operates entirely in mask space and its cost is independent of channel dimensions, making its overhead negligible relative to the compute stage.

[TODO] WISEConv overview. Left: input mask M_{in} on the spatial grid. Middle: construction stage scans tiles, emits ordered worklist $\mathbf{q}$. Right: compute stage consumes $\mathbf{q}$, drives implicit GEMM, writes output $\mathbf{y}$ at active positions only.

Fig. 2. WISEConv two-stage pipeline. The construction stage scans the mask tile by tile and emits an ordered worklist. The compute stage drives the dense convolution datapath using only the worklist positions.

B. Worklist Construction

The construction stage produces the worklist $\mathbf{q}$ from the input mask M_{in} by performing a per-tile gather followed by concatenation. It operates on the tile partition $T_1, \dots, T_N$ defined in Eq. (2). For each output position, the stage first evaluates Eq. (3) to determine the output mask M (requiring k^2 boolean reads of M_{in} per position), then gathers the active indices tile by tile.

a) *Per-tile gather.*: For each tile T_n , the output mask M induces a support set of active indices within the tile:

$$\mathcal{I}_n = \{k \in \{1, \dots, t_h t_w\} : M[T_n[k]] = 1\}. \quad (4)$$

The per-tile subsequence is the gather of T_n by its support set:

$$\mathbf{q}_n = (T_n)_{\mathcal{I}_n} = (T_n[k])_{k \in \mathcal{I}_n}, \quad (5)$$

ordered by ascending k (i.e., row-major within the tile), with length $|\mathbf{q}_n| = |\mathcal{I}_n|$.

b) *Worklist concatenation.*: Let $P = (P_1, \dots, P_N)$ be the permutation of $\{1, \dots, N\}$ determined by nondeterministic parallel block scheduling. The worklist is:

$$\mathbf{q} = (\mathbf{q}_{P_1}, \mathbf{q}_{P_2}, \dots, \mathbf{q}_{P_N}), \quad (6)$$

with total length $|\mathbf{q}| = \|M\|_0$. Each $\mathbf{q}_{P_n}$ retains its internal row-major ordering; only the inter-tile order varies with P .

c) *Cost analysis.*: The construction stage evaluates Eq. (3) for all $H_{out}W_{out}$ output positions, performing k^2 boolean reads of M_{in} per position, plus $\|M\|_0$ integer writes to form $\mathbf{q}$. Its total work per layer is $O(H_{out}W_{out} \cdot k^2)$, independent of C_{in} and C_{out} .

To estimate the overhead relative to computation, we model per-layer latency as $\hat{t} = \hat{t}_{construct} + \hat{t}_{compute}$, assuming each is proportional to its operation count. Construction performs $H_{out}W_{out} \cdot k^2$ operations; computation performs $\rho^l \cdot H_{out}W_{out} \cdot k^2 \cdot C_{in} \cdot C_{out}$ FMAs. Under this model the overhead ratio is:

$$\hat{t}_{construct} / \hat{t}_{compute} \leq 1 / (\rho^l \cdot C_{in} \cdot C_{out}). \quad (7)$$

For a layer with $C_{in} = C_{out} = 64$ and $\rho^l = 5\%$, this gives $\approx 0.5\%$; at $\rho^l = 1\%$ it is $\approx 2.4\%$. Section ?? confirms that the measured construction overhead is consistent with this bound.

Algorithm 1 Worklist Construction

Require: Tiles $T_1, \dots, T_N$, input mask M_{in}

Ensure: Worklist $\mathbf{q}$, output mask M

```

1: for each tile  $T_n$  in parallel do
2:   for each  $k \in \{1, \dots, t_h t_w\}$  do
3:      $g \leftarrow T_n[k]$ 
4:      $M[g] \leftarrow \text{Eq. (3)}$ 
5:   end for
6:    $\mathcal{I}_n \leftarrow \{k : M[T_n[k]] = 1\}$ 
7:    $\mathbf{q}_n \leftarrow (T_n[k])_{k \in \mathcal{I}_n}$ 
8: end for
9:  $\mathbf{q} \leftarrow (\mathbf{q}_{P_1}, \dots, \mathbf{q}_{P_N})$       { $P$  from block scheduling}

```

C. Worklist-driven Convolution

The compute stage fuses position-level selectivity into the dense pipeline's coordinate-driven execution: it computes Eq. (1) for exactly the $\|M\|_0$ worklist entries while preserving the per-position datapath unchanged, summarized as Algorithm 2. Following standard implicit GEMM, the workload is partitioned into compute tiles: the worklist positions are divided into segments of size B_M indexed by b_m , and the C_{out} output channels into segments of size B_N indexed by b_n ; all compute tiles execute in parallel (line 1). Each compute tile first loads its B_M worklist entries into $\mathbf{q}'$ (line 2); this is the sole deviation from the dense pipeline, which uses implicitly sequential coordinates. After initializing the accumulator (line 3), the tile iterates over the $k^2 C_{in}$ reduction dimension in steps of B_K (lines 4–7): at each step b_k , the kernel uses coordinates in $\mathbf{q}'$ to gather the corresponding input slice from $\mathbf{x}$, loads the matching weight block, and accumulates the partial product. Finally, the tile writes the accumulated result back to $\mathbf{y}$ using the linear indices in $\mathbf{q}'$ as output addresses (line 8).

a) *Locality.*: Each compute tile processes B_M consecutive worklist entries. Because construction compacts active positions per tile, consecutive entries in each $\mathbf{q}_{P_n}$ are row-major neighbors in G_{out} whose receptive fields overlap in

Algorithm 2 Worklist-driven Convolution**Require:** Worklist $\mathbf{q}$, input $\mathbf{x}$, weight $\hat{\mathbf{W}}$ **Ensure:** Output $\mathbf{y}$ at positions in $\mathbf{q}$

```

1: for each compute tile  $(b_m, b_n)$  in parallel do
2:    $\mathbf{q}' \leftarrow$  load  $B_M$  entries from  $\mathbf{q}$  for segment  $b_m$ 
3:   Initialize accumulator  $\mathbf{Y} \leftarrow \mathbf{0}_{B_M \times B_N}$ 
4:   for  $b_k = 1, \dots, \lceil k^2 C_{in}/B_K \rceil$  do
5:      $\mathbf{X} \leftarrow$  gather input block from  $\mathbf{x}$  at positions  $\mathcal{N}(\mathbf{q}')$ ,
       reduction step  $b_k$ 
6:      $\hat{\mathbf{W}}' \leftarrow$  load weight block for step  $b_k$ , channels  $b_n$ 
7:      $\mathbf{Y} \leftarrow \mathbf{Y} + \mathbf{X} \cdot \hat{\mathbf{W}}'$ 
8:   end for
9:   Write  $\mathbf{Y}$  to  $\mathbf{y}$  at addresses  $\mathbf{q}'$ 
10: end for

```

contiguous memory. A compute chunk of B_M entries spans at most $\lceil B_M / \min_n |\mathbf{q}_{P_n}| \rceil$ construction tiles regardless of the permutation P ; the nondeterministic inter-tile ordering does not degrade spatial locality relative to the dense raster order.

IV. IMPLEMENTATION

WISEConv is implemented as a CUDA C++ extension exposed to PyTorch via Pybind11, serving as a drop-in replacement for standard convolution layers. All tensors use the channels-last (NHWC) memory layout, ensuring that each active position's gather reads a contiguous memory segment. The construction kernel maps one thread per output position; each thread tests the receptive field against M_{in} with early exit. Within each warp, `__ballot_sync` produces a bitmask of active lanes and `__popc` on masked prefixes yields each active lane's rank, giving the support set $\mathcal{I}_n$ (the ordered indices of active positions) without a sort or prefix-sum pass. One warp-elected thread performs a single `atomicAdd` to reserve a contiguous worklist segment; each active lane writes its index at the reserved base plus its rank. The compute kernel follows a Cutlass-style [?] implicit-GEMM structure with double-buffered shared memory and asynchronous global-to-shared loads (`cp.async` on Ampere+). Tile configurations are selected via Triton-style [?] offline autotuning.

ACKNOWLEDGMENTS

REFERENCES

- [1] D. Wofk, F. Ma, T. Yang, S. Karaman, and V. Sze, "Fastdepth: Fast monocular depth estimation on embedded systems," in *International Conference on Robotics and Automation, ICRA 2019, Montreal, QC, Canada, May 20-24, 2019*. IEEE, 2019, pp. 6101–6108. [Online]. Available: <https://doi.org/10.1109/ICRA.2019.8794182>
- [2] R. J. Bouwmeester, F. Paredes-Vallés, and G. C. H. E. de Croon, "Nanoflownet: Real-time dense optical flow on a nano quadcopter," in *IEEE International Conference on Robotics and Automation, ICRA 2023, London, UK, May 29 - June 2, 2023*. IEEE, 2023, pp. 1996–2003. [Online]. Available: <https://doi.org/10.1109/ICRA48891.2023.10161258>
- [3] M. Sundermeyer, A. Mousavian, R. Triebel, and D. Fox, "Contact-graspnet: Efficient 6-dof grasp generation in cluttered scenes," in *IEEE International Conference on Robotics and Automation, ICRA 2021, Xi'an, China, May 30 - June 5, 2021*. IEEE, 2021, pp. 13 438–13 444. [Online]. Available: <https://doi.org/10.1109/ICRA48506.2021.9561877>
- [4] S. Durvasula, Y. Guan, and N. Vijaykumar, "Ev-conv: Fast CNN inference on event camera inputs for high-speed robot perception," *IEEE Robotics Autom. Lett.*, vol. 8, no. 6, pp. 3174–3181, 2023. [Online]. Available: <https://doi.org/10.1109/LRA.2023.3259731>

- [5] N. Messikommer, D. Gehrig, A. Loquercio, and D. Scaramuzza, "Event-based asynchronous sparse convolutional networks," in *Computer Vision - ECCV 2020 - 16th European Conference, Glasgow, UK, August 23-28, 2020, Proceedings, Part VIII*, ser. Lecture Notes in Computer Science, A. Vedaldi, H. Bischof, T. Brox, and J. Frahm, Eds., vol. 12353. Springer, 2020, pp. 415–431. [Online]. Available: https://doi.org/10.1007/978-3-030-58598-3_25
- [6] M. Parger, C. Tang, C. D. Twigg, C. Keskin, R. Wang, and M. Steinberger, "Deltacnn: End-to-end CNN inference of sparse frame differences in videos," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2022, New Orleans, LA, USA, June 18-24, 2022*. IEEE, 2022, pp. 12 487–12 496. [Online]. Available: <https://doi.org/10.1109/CVPR52688.2022.01217>
- [7] F. Li, G. Li, X. He, and J. Cheng, "Dynamic dual gating neural networks," in *2021 IEEE/CVF International Conference on Computer Vision, ICCV 2021, Montreal, QC, Canada, October 10-17, 2021*. IEEE, 2021, pp. 5310–5319. [Online]. Available: <https://doi.org/10.1109/ICCV48922.2021.00528>
- [8] M. Ren, A. Pokrovsky, B. Yang, and R. Urtasun, "Sbnet: Sparse blocks network for fast inference," in *2018 IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2018, Salt Lake City, UT, USA, June 18-22, 2018*. Computer Vision Foundation / IEEE Computer Society, 2018, pp. 8711–8720. [Online]. Available: http://openaccess.thecvf.com/content_cvpr_2018/html/Ren_SBNNet_Sparse_Blocks_CVPR_2018_paper.html
- [9] A. Habibian, D. Abati, T. S. Cohen, and B. E. Bejnordi, "Skip-convolutions for efficient video processing," in *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2021, virtual, June 19-25, 2021*. Computer Vision Foundation / IEEE, 2021, pp. 2695–2704. [Online]. Available: https://openaccess.thecvf.com/content/CVPR2021/html/Habibian_Skip-Convolution_for_Efficient_Video_Processing_CVPR_2021_paper.html
- [10] Z. Xie, Z. Zhang, X. Zhu, G. Huang, and S. Lin, "Spatially adaptive inference with stochastic feature sampling and interpolation," in *Computer Vision - ECCV 2020 - 16th European Conference, Glasgow, UK, August 23-28, 2020, Proceedings, Part I*, ser. Lecture Notes in Computer Science, A. Vedaldi, H. Bischof, T. Brox, and J. Frahm, Eds., vol. 12346. Springer, 2020, pp. 531–548. [Online]. Available: https://doi.org/10.1007/978-3-030-58452-8_31
- [11] L. Cavigelli, P. Degen, and L. Benini, "Cbinfer: Change-based inference for convolutional neural networks on video data," in *Proceedings of the 11th International Conference on Distributed Smart Cameras, Stanford, CA, USA, September 5-7, 2017*, M. O. Arias-Estrada, C. Micheloni, H. K. Aghajan, O. I. Camps, and V. M. Brea, Eds. ACM, 2017, pp. 1–8. [Online]. Available: <https://doi.org/10.1145/3131885.3131906>
- [12] Z. Teed and J. Deng, "RAFT: recurrent all-pairs field transforms for optical flow," in *Computer Vision - ECCV 2020 - 16th European Conference, Glasgow, UK, August 23-28, 2020, Proceedings, Part II*, ser. Lecture Notes in Computer Science, A. Vedaldi, H. Bischof, T. Brox, and J. Frahm, Eds., vol. 12347. Springer, 2020, pp. 402–419. [Online]. Available: https://doi.org/10.1007/978-3-030-58536-5_24
- [13] C. Wang, A. Bochkovskiy, and H. M. Liao, "Yolov7: Trainable bag-of-freebies sets new state-of-the-art for real-time object detectors," in *IEEE/CVF Conference on Computer Vision and Pattern Recognition, CVPR 2023, Vancouver, BC, Canada, June 17-24, 2023*. IEEE, 2023, pp. 7464–7475. [Online]. Available: <https://doi.org/10.1109/CVPR52729.2023.00721>
- [14] C. Yu, J. Wang, C. Peng, C. Gao, G. Yu, and N. Sang, "Bisenet: Bilateral segmentation network for real-time semantic segmentation," in *Computer Vision - ECCV 2018 - 15th European Conference, Munich, Germany, September 8-14, 2018, Proceedings, Part XIII*, ser. Lecture Notes in Computer Science, V. Ferrari, M. Hebert, C. Sminchisescu, and Y. Weiss, Eds., vol. 11217. Springer, 2018, pp. 334–349. [Online]. Available: https://doi.org/10.1007/978-3-030-01261-8_20
- [15] C. Wang, D. Xu, Y. Zhu, R. M. Martin, C. Lu, L. Fei-Fei, and S. Savarese, "Densefusion: 6d object pose estimation by iterative dense fusion," in *IEEE Conference on Computer Vision and Pattern Recognition, CVPR 2019, Long Beach, CA, USA, June 16-20, 2019*. Computer Vision Foundation / IEEE, 2019, pp. 3343–3352. [Online]. Available: http://openaccess.thecvf.com/content_CVPR_2019/html/Wang_DenseFusion_6D_Object_Pose_Estimation_by_Iterative_Dense_Fusion_CVPR_2019_paper.html
- [16] A. Dosovitskiy, L. Beyer, A. Kolesnikov, D. Weissenborn, X. Zhai, T. Unterthiner, M. Dehghani, M. Minderer, G. Heigold, S. Gelly, J. Uszkoreit, and N. Houlsby, "An image is worth 16x16 words: Transformers for image recognition at scale," in *9th International Conference on Learning Representations, ICLR 2021, Virtual Event,*

- Austria, May 3-7, 2021*. OpenReview.net, 2021. [Online]. Available: <https://openreview.net/forum?id=YicbFdNTTy>
- [17] S. Chetlur, C. Woolley, P. Vandermersch, J. Cohen, J. Tran, B. Catanzaro, and E. Shelhamer, “cudnn: Efficient primitives for deep learning,” *CoRR*, vol. abs/1410.0759, 2014. [Online]. Available: <http://arxiv.org/abs/1410.0759>
 - [18] G. Gallego, T. Delbrück, G. Orchard, C. Bartolozzi, B. Taba, A. Censi, S. Leutenegger, A. J. Davison, J. Conradt, K. Daniilidis, and D. Scaramuzza, “Event-based vision: A survey,” *IEEE Trans. Pattern Anal. Mach. Intell.*, vol. 44, no. 1, pp. 154–180, 2022. [Online]. Available: <https://doi.org/10.1109/TPAMI.2020.3008413>
 - [19] Y. Yan, Y. Mao, and B. Li, “SECOND: sparsely embedded convolutional detection,” *Sensors*, vol. 18, no. 10, p. 3337, 2018. [Online]. Available: <https://doi.org/10.3390/s18103337>
 - [20] H. Tang, S. Yang, Z. Liu, K. Hong, Z. Yu, X. Li, G. Dai, Y. Wang, and S. Han, “Torchsparse++: Efficient training and inference framework for sparse convolution on gpus,” in *Proceedings of the 56th Annual IEEE/ACM International Symposium on Microarchitecture, MICRO 2023, Toronto, ON, Canada, 28 October 2023 - 1 November 2023*. ACM, 2023, pp. 225–239. [Online]. Available: <https://doi.org/10.1145/3613424.3614303>